

LOAN DATASET

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt, seaborn as sns
```

```
In [2]: df=pd.read_csv(r"C:\Users\91756\Documents\python\loan1.csv")
df
```

```
Out[2]:
```

	Home Owner	Marital Status	Annual Income	Defaulted Borrower
0	Yes	Single	125	No
1	No	Married	100	No
2	No	Single	70	No
3	Yes	Married	120	No
4	No	Divorced	95	Yes
5	No	Married	60	No
6	Yes	Divorced	220	No
7	No	Single	85	Yes
8	No	Married	75	No
9	No	Single	90	Yes

```
In [3]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 4 columns):
 #   Column              Non-Null Count  Dtype  
---  --
 0   Home Owner          10 non-null     object  
 1   Marital Status      10 non-null     object  
 2   Annual Income       10 non-null     int64   
 3   Defaulted Borrower  10 non-null     object  
dtypes: int64(1), object(3)
memory usage: 448.0+ bytes
```

```
In [4]: x=df.drop('Defaulted Borrower',axis=1)
y=df['Defaulted Borrower']
```

```
In [5]: df['Marital Status'].value_counts()
```

```
Out[5]: Marital Status
Single      4
Married     4
Divorced    2
Name: count, dtype: int64
```

```
In [6]: H0={"Home Owner":{"Yes":1,"No":0}}
df=df.replace(H0)
print(df)
```

	Home Owner	Marital Status	Annual Income	Defaulted Borrower
0	1	Single	125	No
1	0	Married	100	No
2	0	Single	70	No
3	1	Married	120	No
4	0	Divorced	95	Yes
5	0	Married	60	No
6	1	Divorced	220	No
7	0	Single	85	Yes
8	0	Married	75	No
9	0	Single	90	Yes

```
In [7]: MS={"Marital Status":{'Single':1,'Married':2,'Divorced':3}}
df=df.replace(MS)
print(df)
```

	Home Owner	Marital Status	Annual Income	Defaulted	Borrower
0	1	1	125		No
1	0	2	100		No
2	0	1	70		No
3	1	2	120		No
4	0	3	95		Yes
5	0	2	60		No
6	1	3	220		No
7	0	1	85		Yes
8	0	2	75		No
9	0	1	90		Yes

```
In [8]: x=df.drop('Defaulted Borrower',axis=1)
y=df['Defaulted Borrower']
```

```
In [9]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.7,random_state=42)
x_train.shape,x_test.shape
```

```
Out[9]: ((7, 3), (3, 3))
```

```
In [10]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[10]: ▾ RandomForestClassifier
RandomForestClassifier()
```

```
In [11]: rf=RandomForestClassifier()
```

```
In [12]: params={'max_depth':[2,3,5,10,20],
'min_samples_leaf':[5,10,20,50,100,200],
'n_estimators':[10,25,30,50,100,200]}
```

```
In [13]: from sklearn.model_selection import GridSearchCV
grid_search=GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

```
Out[13]: ▸ GridSearchCV
▸ estimator: RandomForestClassifier
▸ RandomForestClassifier
```

```
In [14]: grid_search.best_score_
```

```
Out[14]: 0.5833333333333333
```

```
In [15]: rf_best=grid_search.best_estimator_
```


In [20]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0   battery_power    2000 non-null   int64
1   blue             2000 non-null   int64
2   clock_speed      2000 non-null   float64
3   dual_sim         2000 non-null   int64
4   fc               2000 non-null   int64
5   four_g           2000 non-null   int64
6   int_memory       2000 non-null   int64
7   m_dep            2000 non-null   float64
8   mobile_wt        2000 non-null   int64
9   n_cores          2000 non-null   int64
10  pc               2000 non-null   int64
11  px_height        2000 non-null   int64
12  px_width         2000 non-null   int64
13  ram              2000 non-null   int64
14  sc_h             2000 non-null   int64
15  sc_w             2000 non-null   int64
16  talk_time        2000 non-null   int64
17  three_g          2000 non-null   int64
18  touch_screen     2000 non-null   int64
19  wifi             2000 non-null   int64
20  price_range      2000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 328.2 KB
```

```
In [21]: x=df.drop('price_range',axis=1)
         y=df['price_range']
```

```
In [22]: from sklearn.model_selection import train_test_split
         x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.4,random_state=42)
         x_train.shape,x_test.shape
```

Out[22]: ((1200, 20), (800, 20))

```
In [23]: from sklearn.ensemble import RandomForestClassifier
         rfc=RandomForestClassifier()
         rfc.fit(x_train,y_train)
```

```
Out[23]: ▾ RandomForestClassifier
         RandomForestClassifier()
```

In [24]: rf=RandomForestClassifier()

```
In [25]: params={'max_depth':[10,20,30,40],
               'min_samples_leaf':[30,23,45,14],
               'n_estimators':[45,34,32,12]}
```

```
In [26]: from sklearn.model_selection import GridSearchCV
         grid_search=GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring="accuracy")
         grid_search.fit(x_train,y_train)
```

```
Out[26]: ▸ GridSearchCV
         ▸ estimator: RandomForestClassifier
           ▸ RandomForestClassifier
```

```
In [27]: grid_search.best_score_
```

Out[27]: 0.8

```
In [28]: rf_best=grid_search.best_estimator_
print(rf_best)
```

RandomForestClassifier(max_depth=10, min_samples_leaf=14, n_estimators=12)

```
In [29]: from sklearn.tree import plot_tree
from sklearn.tree import DecisionTreeClassifier
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[8],feature_names=x.columns,class_names=['LOW','HIGH','MEDIUM','VERYHIGH'],filled=True)

class_names = ['LOW', 'HIGH', 'MEDIUM', 'VERYHIGH']
Text(0.8727272727272727, 0.4375, 'ram <= 2398.0\ngini = 0.696\nsamples = 92\nvalue = [16, 49, 26, 60]\nclass = VERYHIGH'),
Text(0.8363636363636363, 0.3125, 'blue <= 0.5\ngini = 0.603\nsamples = 54\nvalue = [16, 49, 22, 2]\nclass = HIGH'),
Text(0.8181818181818182, 0.1875, 'ram <= 1254.5\ngini = 0.515\nsamples = 29\nvalue = [8, 28, 7, 0]\nclass = HIGH'),
Text(0.8, 0.0625, 'gini = 0.488\nsamples = 14\nvalue = [8, 11, 0, 0]\nclass = HIGH'),
Text(0.8363636363636363, 0.0625, 'gini = 0.413\nsamples = 15\nvalue = [0, 17, 7, 0]\nclass = HIGH'),
Text(0.8545454545454545, 0.1875, 'gini = 0.653\nsamples = 25\nvalue = [8, 21, 15, 2]\nclass = HIGH'),
Text(0.9090909090909091, 0.3125, 'mobile_wt <= 105.0\ngini = 0.121\nsamples = 38\nvalue = [0, 0, 4, 58]\nclass = VERYHIGH'),
Text(0.8909090909090909, 0.1875, 'gini = 0.278\nsamples = 14\nvalue = [0, 0, 4, 20]\nclass = VERYHIGH'),
Text(0.9272727272727272, 0.1875, 'gini = 0.0\nsamples = 24\nvalue = [0, 0, 0, 38]\nclass = VERYHIGH'),
Text(0.9636363636363636, 0.4375, 'px_height <= 685.5\ngini = 0.695\nsamples = 38\nvalue = [14, 19, 18, 3]\nclass = HIGH'),
Text(0.9454545454545454, 0.3125, 'gini = 0.703\nsamples = 15\nvalue = [8, 5, 8, 2]\nclass = LOW'),
Text(0.9818181818181818, 0.3125, 'gini = 0.653\nsamples = 23\nvalue = [6, 14, 10, 1]\nclass = HIGH')

<Figure: 1000x1000px>
```

```
In [31]: imp_df=pd.DataFrame({"varname":x_train.columns,"Imp":rf_best.feature_importances_})
imp_df.sort_values(by="Imp",ascending=False)
```

Out[31]:

	varname	Imp
13	ram	0.707798
0	battery_power	0.074359
12	px_width	0.060707
11	px_height	0.029782
8	mobile_wt	0.018743
6	int_memory	0.016066
16	talk_time	0.013612
2	clock_speed	0.012709
7	m_dep	0.011171
4	fc	0.011156
10	pc	0.010022
14	sc_h	0.008385
15	sc_w	0.008191
9	n_cores	0.008148
3	dual_sim	0.002303
1	blue	0.001899
18	touch_screen	0.001677
5	four_g	0.001606
19	wifi	0.000925
17	three_g	0.000741

Mobile Price

Test Dataset

```
In [32]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt, seaborn as sns
```

```
In [33]: df=pd.read_csv(r"C:\Users\91756\Documents\python\Mobile_Price_Classification_test.csv")
df
```

Out[33]:

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	...	pc	px_height	px_width	ram	sc_h
0	1	1043	1	1.8	1	14	0	5	0.1	193	...	16	226	1412	3476	12
1	2	841	1	0.5	1	4	1	61	0.8	191	...	12	746	857	3895	6
2	3	1807	1	2.8	0	1	0	27	0.9	186	...	4	1270	1366	2396	17
3	4	1546	0	0.5	1	18	1	25	0.5	96	...	20	295	1752	3893	10
4	5	1434	0	1.4	0	11	1	49	0.5	108	...	18	749	810	1773	15
...
995	996	1700	1	1.9	0	0	1	54	0.5	170	...	17	644	913	2121	14
996	997	609	0	1.8	1	0	0	13	0.9	186	...	2	1152	1632	1933	8
997	998	1185	0	1.4	0	1	1	8	0.5	80	...	12	477	825	1223	5
998	999	1533	1	0.5	1	0	0	50	0.4	171	...	12	38	832	2509	15
999	1000	1270	1	0.5	0	4	1	35	0.1	140	...	19	457	608	2828	9

1000 rows × 21 columns



```
In [34]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   id                     1000 non-null   int64  
1   battery_power          1000 non-null   int64  
2   blue                   1000 non-null   int64  
3   clock_speed            1000 non-null   float64 
4   dual_sim               1000 non-null   int64  
5   fc                     1000 non-null   int64  
6   four_g                 1000 non-null   int64  
7   int_memory             1000 non-null   int64  
8   m_dep                  1000 non-null   float64 
9   mobile_wt              1000 non-null   int64  
10  n_cores                1000 non-null   int64  
11  pc                     1000 non-null   int64  
12  px_height              1000 non-null   int64  
13  px_width               1000 non-null   int64  
14  ram                    1000 non-null   int64  
15  sc_h                   1000 non-null   int64  
16  sc_w                   1000 non-null   int64  
17  talk_time              1000 non-null   int64  
18  three_g                1000 non-null   int64  
19  touch_screen           1000 non-null   int64  
20  wifi                   1000 non-null   int64  
dtypes: float64(2), int64(19)
memory usage: 164.2 KB
```

```
In [35]: x=df.drop('four_g',axis=1)
y=df['four_g']
```

```
In [36]: df['dual_sim'].value_counts()
```

```
Out[36]: dual_sim
1      517
0      483
Name: count, dtype: int64
```

```
In [37]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.4,random_state=42)
x_train.shape,x_test.shape
```

```
Out[37]: ((600, 20), (400, 20))
```

```
In [38]: rf=RandomForestClassifier()
```

```
In [39]: params={'max_depth':[2,30,45,34,45],
'min_samples_leaf':[23,34,45,56],
'n_estimators':[10,23,34,78,89]}
```

```
In [40]: from sklearn.model_selection import GridSearchCV
grid_search=GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring="accuracy")
grid_search.fit(x_test,y_test)
```

```
Out[40]:
> GridSearchCV
> estimator: RandomForestClassifier
  > RandomForestClassifier
```

```
In [41]: grid_search.best_score_
```

```
Out[41]: 0.655
```

```
In [42]: rf_best=grid_search.best_estimator_
print(rf_best)
```

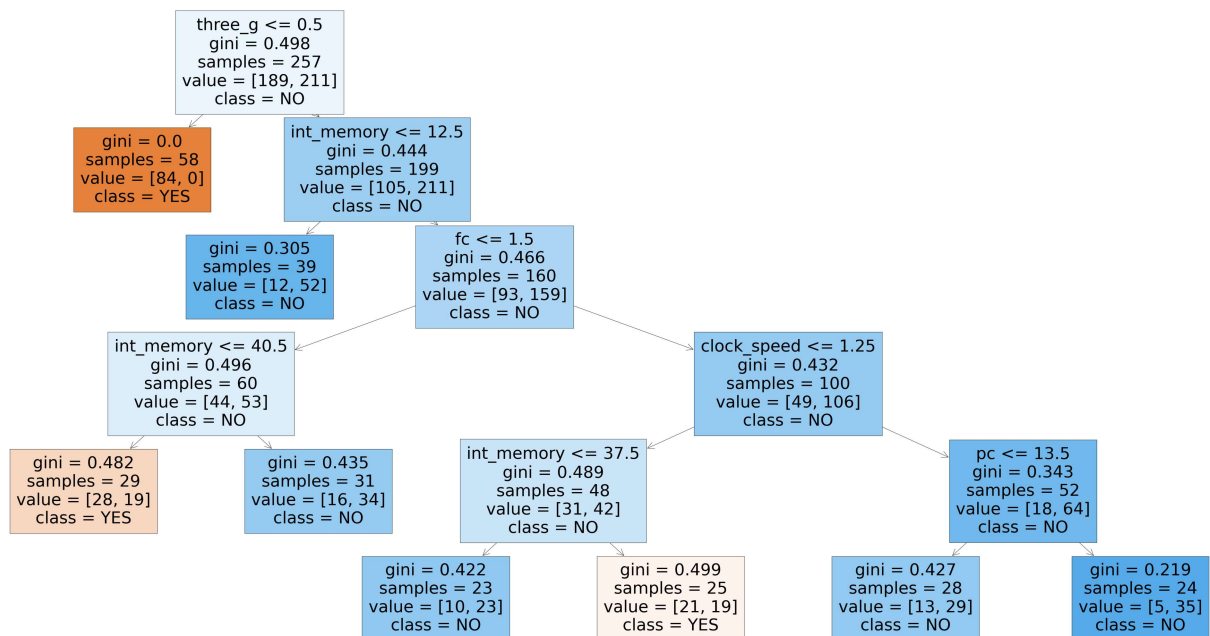
```
RandomForestClassifier(max_depth=30, min_samples_leaf=23, n_estimators=78)
```

```
In [43]: from sklearn.tree import plot_tree
         from sklearn.tree import DecisionTreeClassifier
         plt.figure(figsize=(80,40))
         plot_tree(rf_best.estimators_[6],feature_names=x.columns,class_names=['YES','NO'],filled=True)
```

```

Out[43]: [Text(0.22727272727272727, 0.9166666666666666, 'three_g <= 0.5\nngini = 0.498\nnsamples = 257\nnvalue = [189, 211]\nnclass = NO'),
Text(0.13636363636363635, 0.75, 'gini = 0.0\nnsamples = 58\nnvalue = [84, 0]\nnclass = YES'),
Text(0.3181818181818182, 0.75, 'int_memory <= 12.5\nngini = 0.444\nnsamples = 199\nnvalue = [105, 211]\nnclass = NO'),
Text(0.22727272727272727, 0.5833333333333334, 'gini = 0.305\nnsamples = 39\nnvalue = [12, 52]\nnclass = NO'),
Text(0.4090909090909091, 0.5833333333333334, 'fc <= 1.5\nngini = 0.466\nnsamples = 160\nnvalue = [93, 159]\nnclass = NO'),
Text(0.18181818181818182, 0.4166666666666667, 'int_memory <= 40.5\nngini = 0.496\nnsamples = 60\nnvalue = [44, 53]\nnclass = NO'),
Text(0.09090909090909091, 0.25, 'gini = 0.482\nnsamples = 29\nnvalue = [28, 19]\nnclass = YES'),
Text(0.2727272727272727, 0.25, 'gini = 0.435\nnsamples = 31\nnvalue = [16, 34]\nnclass = NO'),
Text(0.6363636363636364, 0.4166666666666667, 'clock_speed <= 1.25\nngini = 0.432\nnsamples = 100\nnvalue = [49, 106]\nnclass = NO'),
Text(0.45454545454545453, 0.25, 'int_memory <= 37.5\nngini = 0.489\nnsamples = 48\nnvalue = [31, 42]\nnclass = NO'),
Text(0.36363636363636365, 0.08333333333333333, 'gini = 0.422\nnsamples = 23\nnvalue = [10, 23]\nnclass = NO'),
Text(0.5454545454545454, 0.08333333333333333, 'gini = 0.499\nnsamples = 25\nnvalue = [21, 19]\nnclass = YES'),
Text(0.8181818181818182, 0.25, 'pc <= 13.5\nngini = 0.343\nnsamples = 52\nnvalue = [18, 64]\nnclass = NO'),
Text(0.7272727272727273, 0.08333333333333333, 'gini = 0.427\nnsamples = 28\nnvalue = [13, 29]\nnclass = NO'),
Text(0.9090909090909091, 0.08333333333333333, 'gini = 0.219\nnsamples = 24\nnvalue = [5, 35]\nnclass = NO')]

```



```
In [44]: rf_best.feature_importances_
```

```
Out[44]: array([0.08007951, 0.04861858, 0.00215825, 0.03492141, 0.00484123,
                0.01816655, 0.04442906, 0.02061731, 0.06007986, 0.0449101 ,
                0.02074206, 0.03977116, 0.03887831, 0.08717415, 0.02919748,
                0.01727266, 0.02218767, 0.37199406, 0.00876095, 0.00519963])
```



```
In [45]: imp_df=pd.DataFrame({"varname":x_test.columns,"Imp":rf_best.feature_importances_})
imp_df.sort_values(by="Imp",ascending=False)
```

Out[45]:

	varname	Imp
17	three_g	0.371994
13	ram	0.087174
0	id	0.080080
8	mobile_wt	0.060080
1	battery_power	0.048619
9	n_cores	0.044910
6	int_memory	0.044429
11	px_height	0.039771
12	px_width	0.038878
3	clock_speed	0.034921
14	sc_h	0.029197
16	talk_time	0.022188
10	pc	0.020742
7	m_dep	0.020617
5	fc	0.018167
15	sc_w	0.017273
18	touch_screen	0.008761
19	wifi	0.005200
4	dual_sim	0.004841
2	blue	0.002158

In []: